

The Ultimate Kubernetes Handbook

With Deeply Explained Code Comments

Your Personal AI Assistant

December 14, 2025

Contents

1 Core Concepts: The Foundation	2
1.1 Namespace (Virtual Clusters)	2
1.2 Pods (The Atomic Unit)	2
1.3 Deployment (The Manager)	2
2 Networking: Connecting Apps	3
2.1 Services (Internal & External Access)	3
2.2 Ingress (The Smart Router)	3
3 Storage: Persistent Data	3
3.1 PersistentVolume (The Physical Disk)	3
3.2 PersistentVolumeClaim (The Ticket)	4
4 Batch Workloads	4
4.1 Job (One-Off Task)	4
4.2 CronJob (Scheduled Task)	4
5 Mini Project: Full Stack Architecture	6
5.1 Step 1: The Database (Redis)	6
5.2 Step 2: The Backend (Python API)	6
5.3 Step 3: The Frontend (React)	7

1 Core Concepts: The Foundation

1.1 Namespace (Virtual Clusters)

```

1 apiVersion: v1
2 kind: Namespace
3 metadata:
4   # 1. THE FOLDER NAME
5   # Think of this as a folder on your computer.
6   # Resources inside here are isolated from resources in other folders.
7   # You delete the Namespace -> You delete everything inside it.
8   name: production-env

```

Listing 1: Namespace Configuration

1.2 Pods (The Atomic Unit)

```

1 apiVersion: v1
2 kind: Pod
3 metadata:
4   # 1. THE POD'S ID CARD
5   # Unique name for this specific object.
6   # Used in commands like: 'kubectl delete pod my-first-pod'
7   name: my-first-pod
8
9   # 2. THE LOCATION
10  # Places this Pod inside the folder we created above.
11  namespace: production-env
12
13   # 3. THE STICKY NOTE (Categorization)
14   # Services use this label to find this Pod.
15   # It is arbitrary: you could say 'tier: backend' or 'fruit: banana'.
16   labels:
17     app: demo-app
18
19   spec:
20     containers:
21       # 4. THE INTERNAL NICKNAME
22       # Only used for logs: 'kubectl logs my-first-pod -c nginx-container'
23       - name: nginx-container
24
25       # 5. THE SOFTWARE PACKAGE
26       # Downloads 'nginx' version 'alpine' from Docker Hub.
27       image: nginx:alpine
28
29       ports:
30         # 6. THE EXPOSED PORT
31         # This is purely informational. It tells humans "I listen on port 80".
32         - containerPort: 80

```

Listing 2: Pod Configuration

1.3 Deployment (The Manager)

```

1 apiVersion: apps/v1
2 kind: Deployment
3 metadata:
4   # 1. THE MANAGER'S NAME
5   # This is the name of the Deployment object itself.
6   name: web-deployment
7   namespace: production-env
8
9   spec:
10    # 2. THE DESIRED TEAM SIZE
11    # "I want exactly 3 clones of this pod running at all times."
12    # If one dies, I will start a new one to keep this number at 3.
13    replicas: 3
14
15    # 3. THE SEARCH FILTER (Crucial!)

```

```

15 # The Manager asks: "Which pods belong to me?"
16 # Answer: "Any pod wearing the 'app: demo-app' badge."
17 selector:
18   matchLabels:
19     app: demo-app # <-- MUST MATCH Point #5 below
20
21 # 4. THE BLUEPRINT (Cookie Cutter)
22 # If I need to create a new Pod, I will use this template.
23 template:
24   metadata:
25     # 5. THE BADGE
26     # I will stick this label on every new Pod I create.
27     # This ensures the Selector (Point #3) can find them.
28   labels:
29     app: demo-app
30 spec:
31   containers:
32     - name: nginx
33       image: nginx:1.14.2

```

Listing 3: Deployment Configuration

2 Networking

2.1 Services (Internal Access)

```

1 apiVersion: v1
2 kind: Service
3 metadata:
4   # 1. THE DNS NAME
5   # Other pods can talk to this service by using this name.
6   # Example URL: "http://web-service"
7   name: web-service
8   namespace: production-env
9 spec:
10  # 2. THE TYPE OF ACCESS
11  # ClusterIP = Internal Only (Default).
12  # NodePort = Open a port on the server firewall.
13  # LoadBalancer = Ask AWS/GCP for a Public IP.
14  type: ClusterIP
15
16  # 3. THE TARGET FINDER
17  # "Send traffic to any Pod with this label."
18  selector:
19    app: demo-app
20
21 ports:
22   - protocol: TCP
23     # 4. THE FRONT DOOR (Service Port)
24     # This is the port other apps use to call the Service.
25     port: 80
26
27     # 5. THE BACK DOOR (Container Port)
28     # This must match the port your container is actually listening on.
29     targetPort: 80

```

Listing 4: Service Configuration

2.2 Ingress (External Routing)

```

1 apiVersion: networking.k8s.io/v1
2 kind: Ingress
3 metadata:
4   name: main-ingress
5   namespace: production-env
6   annotations:
7     # 1. THE CONTROLLER INSTRUCTION

```

```

8   # Specific configuration for Nginx (e.g., strip the path).
9   nginx.ingress.kubernetes.io/rewrite-target: /
10  spec:
11    # 2. THE POLICE OFFICER
12    # Tells the cluster which Ingress Controller should enforce these rules.
13    ingressClassName: nginx
14
15  rules:
16    # 3. THE DOMAIN NAME
17    # "I only care about traffic for myapp.com"
18    - host: myapp.com
19      http:
20        paths:
21          # 4. THE PATH RULE
22          # "If the user visits myapp.com/api..."
23          - path: /api
24            pathType: Prefix
25            backend:
26              service:
27                # 5. THE DESTINATION
28                # "...send them to the backend-service on port 5000."
29                name: backend-service
30                port: { number: 5000 }

```

Listing 5: Ingress Rules

3 Storage

3.1 PersistentVolume (The Hardware)

```

1  apiVersion: v1
2  kind: PersistentVolume
3  metadata:
4    # 1. THE HARD DRIVE NAME
5    # Managed by the IT Admin, not the Developer.
6    name: pv-10gb
7  spec:
8    # 2. THE SIZE
9    # How big is this physical disk?
10   capacity:
11     storage: 10Gi
12
13   # 3. THE ACCESS RULES
14   # ReadWriteOnce = Only one node can mount this at a time (like a USB stick).
15   # ReadWriteMany = Many nodes can share it (like a Shared Folder).
16   accessModes:
17     - ReadWriteOnce
18
19   # 4. THE DRIVER
20   # "hostPath" uses the Node's local disk (Testing only).
21   # In production, use 'awsElasticBlockStore' or 'gcePersistentDisk'.
22   hostPath:
23     path: "/mnt/data"

```

Listing 6: PV Definition

3.2 PersistentVolumeClaim (The Ticket)

```

1  apiVersion: v1
2  kind: PersistentVolumeClaim
3  metadata:
4    # 1. THE TICKET NAME
5    # The Pod will use this name to claim storage.
6    name: db-storage-claim
7    namespace: production-env
8  spec:
9    # 2. THE REQUEST

```

```

10 # "I need a drive that supports Single-Node writing."
11 accessModes:
12   - ReadWriteOnce
13 resources:
14   requests:
15     # 3. THE SIZE NEEDED
16     # "Find me a PV that has at least 5GB free."
17     storage: 5Gi

```

Listing 7: PVC Request

4 Batch Jobs

4.1 CronJob (Scheduled)

```

1 apiVersion: batch/v1
2 kind: CronJob
3 metadata:
4   name: nightly-backup
5 spec:
6   # 1. THE SCHEDULE
7   # Minute / Hour / Day-Month / Month / Day-Week
8   # "Run at 00:00 (Midnight) every day."
9   schedule: "0 0 * * *"
10
11   # 2. THE JOB TEMPLATE
12   # Every midnight, create a Job object with these settings.
13   jobTemplate:
14     spec:
15       # 3. THE POD TEMPLATE
16       # The Job will create a Pod with these settings.
17       template:
18         spec:
19           containers:
20             - name: backup
21               image: backup-tool
22
23             # 4. THE RESTART POLICY
24             # "If the script crashes, try again (OnFailure)."
25             # "Do not restart endlessly like a web server."
26             restartPolicy: OnFailure

```

Listing 8: CronJob Example

5 Mini Project: Full Stack

5.1 Backend (Python API)

```

1 apiVersion: apps/v1
2 kind: Deployment
3 metadata:
4   name: backend-api
5 spec:
6   replicas: 2
7   selector: { matchLabels: { app: backend } }
8   template:
9     metadata: { labels: { app: backend } }
10    spec:
11      containers:
12        - name: python-api
13          image: my-python-api:v1
14
15        # 1. CONNECTING TO REDIS
16        # We don't use IP addresses. We use the Service Name.
17        env:
18          - name: REDIS_HOST

```

```

19         # This name 'redis-service' matches the Service YAML below.
20         value: "redis-service"
21
22     apiVersion: v1
23     kind: Service
24     metadata:
25       name: backend-service
26     spec:
27       # 1. INTERNAL ONLY
28       # The public cannot access the API directly. Only the Frontend can.
29       type: ClusterIP
30       selector: { app: backend }
31       ports: [{ port: 5000, targetPort: 5000 }]

```

Listing 9: Backend Deployment

5.2 Frontend (React)

```

1  apiVersion: apps/v1
2  kind: Deployment
3  metadata:
4    name: frontend-ui
5  spec:
6    replicas: 3
7    selector: { matchLabels: { app: frontend } }
8    template:
9      metadata: { labels: { app: frontend } }
10     spec:
11       containers:
12         - name: react-ui
13           image: my-react-app:v1
14
15         # 1. CONNECTING TO BACKEND
16         # The frontend browser needs to know where to send API requests.
17         env:
18           - name: API_URL
19             value: "http://backend-service:5000"
20
21     apiVersion: v1
22     kind: Service
23     metadata:
24       name: public-frontend
25     spec:
26       # 1. PUBLIC ACCESS
27       # This asks the Cloud Provider (AWS/GCP) for a Real Public IP.
28       type: LoadBalancer
29       selector: { app: frontend }
30       ports:
31         - port: 80          # Users hit port 80 (HTTP)
32           targetPort: 80   # Container listens on 80

```

Listing 10: Frontend Public Access